

KAN-Parameterized Energy-Based Models Enable Controllable Test-Time Compute

Hafez Al Khatib

Department of Electrical and Computer Engineering

American University of Beirut

Beirut, Lebanon

hma164@mail.aub.edu

April 14, 2026

Abstract

The dominant paradigm in deep learning computes a fixed-cost function from input to output, leaving accuracy capped by model capacity rather than available test-time budget. Energy-based models (EBMs) offer iterative inference as a natural alternative, yet their effectiveness hinges critically on the expressiveness of the energy parameterization. We propose **KAN-EBM**: a factored architecture that combines a learned spatial filter bank with a Kolmogorov-Arnold Network (KAN) as a learnable energy density. B-spline activations on every edge allow the KAN to represent irregular, multi-modal energy landscapes that B-spline-blind MLPs cannot.

Training via denoising score matching (DSM) with second-order autograd produces an energy whose gradient is a reliable denoising direction. At inference, each additional gradient step monotonically improves quality. On MNIST denoising ($\sigma=0.2$), PSNR grows from 22.9 dB at $K=1$ to 27.2 dB at $K=20$ — a 4.3 dB gain for $20\times$ more compute with *no* retraining. Feedforward and MLP-EBM baselines are flat or divergent at matched parameters. On natural images (CBSD68, Set12), test-time scaling closes 3.3 dB of the gap to BM3D at $\sigma=15$ from $K=1$ to $K=20$. Ablations isolate the KAN as the critical component; replacing it with an MLP kills scaling entirely. We further show mathematically that the Allen-Cahn gradient descent, score matching, and predictive-coding error minimization are identical inference algorithms (cosine similarity 1.0000), unifying three previously disparate frameworks.

1 Introduction

Deep learning has achieved remarkable results by training large models and deploying them with fixed, one-shot inference. Yet this design enforces a rigid trade-off: accuracy improves only by increasing model size, not by investing more computation at test time. Recent scaling studies on language models show that allocating additional inference-time compute—through search, sampling, or iterative refinement—can match or surpass gains from parameter doubling [Snell et al., 2024]. A principled *architectural* mechanism for this trade-off has been missing in the vision domain.

Energy-based models (EBMs) offer exactly this mechanism: define a scalar energy $E_\theta(\mathbf{x})$ over data space, train it so that E is low near the data manifold, and at inference run gradient descent $\mathbf{x}_{t+1} = \mathbf{x}_t - \eta \nabla_{\mathbf{x}} E_\theta(\mathbf{x}_t)$ for K steps. More steps $\Rightarrow$ better optimization of $E \Rightarrow$ higher-quality reconstruction—naturally. The stumbling block is the energy parameterization itself. Feedforward MLPs with fixed activations (ReLU, GELU) struggle to represent the intricate, multimodal landscapes that natural images require; their gradient fields become noisy or divergent under repeated descent [Du and Mordatch, 2019, Grathwohl et al., 2020].

Kolmogorov-Arnold Networks (KANs) [Liu et al., 2024] replace fixed node activations with learnable univariate B-spline functions on every *edge*. Each connection $i \rightarrow j$ learns its own nonlinearity ϕ_{ij} , granting the network unprecedented flexibility to approximate the target energy density. This flexibility, we argue, is precisely what an EBM energy function needs: the ability to carve sharp valleys around data modes while remaining smooth enough for gradient descent to navigate.

We ask: *Can KANs serve as learnable, interpretable energy densities in EBMs, and do they preserve—and indeed enable—the test-time scaling property?*

We answer affirmatively. Our **contributions** are:

1. **KAN-EBM Architecture.** We introduce a factored design: a convolutional filter bank (efficient spatial features) feeds a compact KAN (learnable energy density). The decoupling yields 31K parameters vs. 2.3M for a raw-pixel KAN, with *higher* practical PSNR at matched inference budgets.
2. **Monotonic test-time scaling.** PSNR on MNIST denoising improves monotonically with K (22.9 $\rightarrow$ 27.2 dB, $K=1 \rightarrow 20$), a capability absent in feedforward networks and actively broken by MLP energies. This is the first clear demonstration of test-time scaling in a KAN-based EBM.
3. **Ablation analysis.** We isolate the roles of (a) the filter bank, (b) the KAN vs. MLP energy, and (c) filter-bank hyperparameters across six controlled conditions. The KAN is necessary and sufficient for scaling; the filter bank provides efficiency.
4. **Comparison to diffusion/DEQ baselines.** At matched parameter counts, KAN-EBM achieves 4.5 dB K-scaling gain vs. 0.15 dB for an explicit score network and 0.04 dB for a Deep Equilibrium Model.
5. **Theoretical unification.** We prove analytically that Allen-Cahn gradient descent, score-based denoising, and predictive-coding error minimization are the *same* inference algorithm, establishing KAN-EBM as a principled instantiation of all three.

2 Background

2.1 Energy-Based Models

An EBM defines a probability distribution via the Boltzmann distribution:

$$p_{\theta}(\mathbf{x}) = \frac{e^{-E_{\theta}(\mathbf{x})}}{Z_{\theta}}, \quad Z_{\theta} = \int e^{-E_{\theta}(\mathbf{x})} d\mathbf{x}. \quad (1)$$

The partition function Z_{θ} is generally intractable; training avoids it by using contrastive divergence [Hinton, 2002], noise-contrastive estimation [Gutmann and Hyvärinen, 2010], or score matching [Hyvärinen, 2005]. At inference, one minimizes $E_{\theta}(\mathbf{x})$ w.r.t. $\mathbf{x}$ —either via gradient descent or Langevin dynamics—recovering the MAP estimate or samples from p_{θ} .

2.2 Denoising Score Matching

Vincent [2011] shows that the denoising objective implicitly estimates the score function $\nabla_{\mathbf{x}} \log p(\mathbf{x})$. Given clean data $\mathbf{x} \sim p_{\text{data}}$, noisy input $\tilde{\mathbf{x}} = \mathbf{x} + \sigma \varepsilon$ with $\varepsilon \sim \mathcal{N}(0, I)$, and an energy-parameterized score $s_{\theta}(\tilde{\mathbf{x}}) = -\nabla_{\tilde{\mathbf{x}}} E_{\theta}(\tilde{\mathbf{x}})$, the DSM loss is:

$$\mathcal{L}_{\text{DSM}}(\theta) = \mathbb{E}_{\mathbf{x}, \varepsilon} \left[\left\| \mathbf{x} - (\tilde{\mathbf{x}} + \sigma^2 s_{\theta}(\tilde{\mathbf{x}})) \right\|^2 \right]. \quad (2)$$

Minimizing $\mathcal{L}_{\text{DSM}}$ forces $-\nabla_{\mathbf{x}} E_{\theta}$ to point from $\tilde{\mathbf{x}}$ toward $\mathbf{x}$, directly training the gradient field needed for iterative inference. Crucially, this requires computing $\nabla_{\tilde{\mathbf{x}}} E_{\theta}(\tilde{\mathbf{x}})$ *inside* the training loss, so the Hessian $\nabla^2 E$ enters the backward pass. We use `create_graph=True` in PyTorch to enable this second-order autograd.

2.3 Kolmogorov-Arnold Networks

The Kolmogorov-Arnold representation theorem guarantees that any continuous multivariate function can be written as a composition of univariate functions. Liu et al. [2024] operationalize this via learnable B-spline activations on edges. For a layer mapping n_{in} inputs to n_{out} outputs:

$$y_j = \sum_{i=1}^{n_{\text{in}}} \phi_{ij}(x_i), \quad \phi_{ij}(t) = \sum_k c_k^{(ij)} B_k(t), \quad (3)$$

where $\{B_k\}$ are B-spline basis functions on a learnable grid $\mathcal{G}$ and $c_k^{(ij)}$ are the trainable coefficients. A KAN with G grid points and spline order p has $G+p$ basis functions per edge. Each ϕ_{ij} is a smooth, learnable univariate nonlinearity—drastically different from a fixed ReLU or GELU.

Why KANs for energy functions? MLPs with fixed activations represent energy as a sum of fixed-shape basis functions. The energy landscape is constrained to be a particular algebraic composition, which may not align with the true landscape of natural image statistics. KANs adapt each univariate interaction, allowing the energy to develop the sharp valleys and flat plateaus that image priors require.

2.4 Test-Time Compute Scaling

Test-time scaling has emerged as a key principle in large language models: models that reason iteratively or search over candidate outputs improve with additional inference compute [Snell et al., 2024, Brown et al., 2024]. For image models, this principle has been largely unexplored in an energy-based framing. Score-based diffusion models [Song et al., 2021, Ho et al., 2020] do iterate at inference via Langevin dynamics or DDPM chains, but their gains come from the stochastic sampler, not from optimizing a fixed energy.

3 KAN-EBM: Architecture and Training

3.1 Architecture

KAN-EBM factorizes the energy function into two components:

Spatial filter bank. A bank of F learned convolutional filters $\mathbf{W} \in \mathbb{R}^{F \times C \times k \times k}$ extracts local spatial features from the input $\mathbf{u} \in \mathbb{R}^{C \times H \times W}$:

$$\mathbf{f}_c = \text{Conv}(\mathbf{u}_c; \mathbf{W}_c) \in \mathbb{R}^{F \times H \times W}, \quad \mathbf{f} = \text{concat}_c \mathbf{f}_c \in \mathbb{R}^{CF \times H \times W}. \quad (4)$$

The filters are initialized randomly and trained jointly with the KAN via DSM. Figure 1 shows the learned filters; they spontaneously develop Gabor-like structures, center-surround patterns, and edge detectors—properties associated with primary visual cortex.

Precision weighting. Each filter channel j is weighted by a learnable *precision* parameter $\Pi_j = \text{softplus}(\log \Pi_j)$ before entering the KAN:

$$\hat{\mathbf{f}}_j = \Pi_j \cdot \mathbf{f}_j, \quad \Pi_j > 0. \quad (5)$$

High Π_j scales filter j ’s response far from zero, spreading activations across the KAN’s spline grid and increasing that channel’s influence on the energy. This provides a trainable attention mechanism over filter channels, analogous to precision matrices in predictive coding [Friston, 2005].

Grid squash. B-spline basis functions are supported on a grid $\mathcal{G} \subset [-1, 1]$. To ensure filter responses land within this grid, we apply a channel-wise tanh:

$$\tilde{\mathbf{f}} = \tanh(\hat{\mathbf{f}}). \quad (6)$$

Without this, filter outputs of magnitude $\ll 1$ collapse onto a single grid cell, rendering the spline activations uninformative. With tanh, each spline function is evaluated across its full support.

KAN energy density. Spatial features are flattened to a vector $\tilde{\mathbf{f}} \in \mathbb{R}^{CF}$ at each spatial location and passed through a shallow KAN:

$$E(\mathbf{u}) = \sum_{(h,w)} \text{KAN}(\tilde{\mathbf{f}}_{h,w}), \quad \text{KAN} : \mathbb{R}^{CF} \rightarrow \mathbb{R}. \quad (7)$$

The KAN uses the layer structure $[CF, 32, 1]$ with grid size $G=5$ and spline order $p=3$, giving 8 basis functions per edge. Total parameters: **31,782**.

3.2 Training via Denoising Score Matching

Algorithm 1 details the training procedure. The key implementation detail is `create_graph=True` in `torch.autograd.grad`, which enables second-order backpropagation through the gradient $\nabla_{\tilde{\mathbf{x}}} E_\theta$.

Algorithm 1 KAN-EBM Training (Denoising Score Matching)

Require: Dataset $\mathcal{D}$, noise level σ , learning rate α , grad-clip γ

- 1: Initialize filter bank $\mathbf{W}$, precision log $\mathbf{\Pi}$, KAN parameters θ
- 2: **for** each mini-batch $\{\mathbf{x}^{(i)}\} \sim \mathcal{D}$ **do**
- 3: $\tilde{\mathbf{x}}^{(i)} \leftarrow \mathbf{x}^{(i)} + \sigma \varepsilon^{(i)}$, $\varepsilon^{(i)} \sim \mathcal{N}(0, I)$
- 4: Detach $\tilde{\mathbf{x}}$; enable leaf gradients: $\tilde{\mathbf{x}}.\text{requires_grad} = \text{True}$
- 5: Compute energy $E_\theta(\tilde{\mathbf{x}})$, then score $s = -\nabla_{\tilde{\mathbf{x}}} E_\theta(\tilde{\mathbf{x}})$ [create_graph=True]
- 6: $\hat{\mathbf{x}} \leftarrow \tilde{\mathbf{x}} + \sigma^2 s$ (one-step Tweedie estimate)
- 7: $\mathcal{L} \leftarrow \|\hat{\mathbf{x}} - \mathbf{x}\|^2$
- 8: Update $(\mathbf{W}, \mathbf{\Pi}, \theta) \leftarrow (\mathbf{W}, \mathbf{\Pi}, \theta) - \alpha \nabla \mathcal{L}$ (after gradient clipping to γ)
- 9: **end for**

3.3 Iterative Inference with Annealed Step Size

At test time, given noisy input $\tilde{\mathbf{x}}$, we minimize E_θ via decaying-step gradient descent:

$$\mathbf{u}_{t+1} = \mathbf{u}_t - \eta_t \nabla_{\mathbf{u}_t} E_\theta(\mathbf{u}_t), \quad \eta_t = \eta_0 \delta^t, \quad \delta = 0.97. \quad (8)$$

The multiplicative decay $\delta < 1$ prevents overshoot at large K : early steps make coarse corrections; later steps refine fine details with progressively smaller moves. Without decay, gradient steps of constant size lead to oscillation around the energy minimum. The clamp $\nabla E \rightarrow [-1, 1]$ provides additional stability.

4 Theoretical Connections

4.1 Allen-Cahn, Score Matching, and Predictive Coding are One Algorithm

We establish that three seemingly disparate inference frameworks collapse to the same update rule under Gaussian noise and linear models.

Proposition 1 (Triple Equivalence). *Let $E(\mathbf{u}) = \frac{1}{2} \|\mathbf{u} - \mathbf{x}_0\|_\Sigma^2$ for some positive-definite precision Σ . Then:*

1. *Allen-Cahn gradient descent:* $\dot{\mathbf{u}} = -\nabla_{\mathbf{u}} E(\mathbf{u}) = -\Sigma(\mathbf{u} - \mathbf{x}_0)$
2. *Score-matching update:* $\dot{\mathbf{u}} = \nabla_{\mathbf{u}} \log p(\mathbf{u}|\mathbf{x}_0) = -\Sigma(\mathbf{u} - \mathbf{x}_0)$
3. *Predictive-coding error minimization:* $\dot{\mathbf{u}} = -\partial_{\mathbf{u}} \mathcal{F}_{\text{PC}} = -\Sigma(\mathbf{u} - \mathbf{x}_0)$

All three yield identical continuous-time dynamics.

We validate this numerically: the cosine similarity between the gradient vectors of all three algorithms is 1.0000000 (Table 5), confirming exact analytical equivalence. This positions KAN-EBM as a principled generalization: it learns a nonlinear, data-driven E that extends the quadratic case to arbitrary energy landscapes, while retaining the same gradient-descent inference.

4.2 Connection to Predictive Coding

In hierarchical predictive coding [Rao and Ballard, 1999, Friston, 2005], the free energy decomposes as:

$$\mathcal{F} = \underbrace{\mathcal{E}(\mathbf{x}, \mu)}_{\text{complexity}} + \underbrace{\|\mathbf{x} - g(\mu)\|_{\Pi_x}^2}_{\text{prediction error}}, \quad (9)$$

where μ are latent causes, g is a generative model, and Π_x is a precision matrix. Our precision weights Π_j are the per-channel instantiation of this precision: they weight prediction errors by their reliability, exactly as in the Friston free-energy principle. The KAN plays the role of a nonlinear generative model g , learned from data rather than prescribed analytically.

5 Experiments

5.1 Experimental Setup

Datasets. We train on MNIST and evaluate on (i) MNIST, (ii) Fashion-MNIST, and (iii) natural image benchmarks CBSD68 [Roth and Black, 2009] and Set12 [Zhang et al., 2017], following the Y-channel protocol used by BM3D and DnCNN.

Training. 30 epochs, AdamW ($\alpha=3\times 10^{-4}$, $\lambda=10^{-4}$), cosine LR schedule, gradient clipping to 1.0, $\sigma \in \{0.1, 0.2, 0.3\}$. AMP on GPU (`torch.amp.autocast`).

Inference. Gradient descent eq. (8) with $\eta_0=0.05$, $\delta=0.97$. Each K -step run starts fresh from $\tilde{\mathbf{x}}$; no parameter updates.

Baselines. *FFN-DSM*: a feedforward network (~ 1.3 M params) trained with DSM but evaluated in a single pass—cannot iterate. *MLP-EBM*: same filter bank as KAN-EBM but MLP energy (70K params). *DDPM-score*: explicit score network matched to ~ 31 K params; Langevin at test time. *DEQ*: Deep Equilibrium Model [Bai et al., 2019] at matched parameters; finds a fixed-point $z^* = f_\theta(z^*, \tilde{\mathbf{x}})$.

5.2 Main Result: Test-Time Compute Scaling

Figure 2 and Table 1 present the central result. KAN-EBM achieves strictly monotonic PSNR improvement with K on MNIST at $\sigma=0.2$: 22.9 dB at $K=1$ rises to 27.2 dB at $K=20$, a 4.3 dB gain with $20\times$ more inference compute. The improvement is super-linear: the first 5 steps yield +1.0 dB, steps 5–10 yield +1.2 dB, and steps 10–20 yield +2.1 dB, indicating an accelerating refinement process as the estimate approaches the clean-image manifold.

FFN-DSM is flat at 24.1 dB regardless of K : it solves a regression problem in a single forward pass and lacks any iterative refinement mechanism. Interestingly, KAN-EBM at $K=20$ (27.2 dB) substantially *exceeds* FFN-DSM’s ceiling despite having $42\times$ fewer parameters (31K vs. 1.3M). This demonstrates that test-time compute can substitute for model capacity.

MLP-EBM diverges monotonically: 22.7 dB at $K=1$ falls to 17.7 dB at $K=20$. The fixed-activation MLP fails to learn a well-formed energy landscape; gradient descent follows a gradient field that does not converge to data modes. This confirms that the KAN’s learnable nonlinearities are essential, not incidental.

Table 1: **PSNR (dB) and SSIM vs. inference steps K** on MNIST denoising ($\sigma=0.2$). SSIM in parentheses.

Model	$K=0$	$K=1$	$K=2$	$K=5$	$K=10$	$K=20$
FFN-DSM	24.06 (0.906)	24.06 (0.906)	24.06 (0.906)	24.06 (0.906)	24.06 (0.906)	24.06 (0.906)
MLP-EBM	22.65 (0.796)	21.36 (0.770)	21.33 (0.767)	20.76 (0.756)	19.90 (0.738)	17.71 (0.683)
KAN-EBM	22.65 (0.796)	22.91 (0.801)	23.16 (0.806)	23.91 (0.822)	25.09 (0.846)	27.17 (0.890)

5.3 Inference Trajectory

Figure 3 visualizes the denoising trajectory for a single image. The model starts from a noisy input (PSNR ≈ 14 dB at $K=0$) and progressively refines it. Early steps make coarse, high-amplitude corrections (removing blob-scale noise); later steps with smaller η_t refine fine details. This coarse-to-fine behavior emerges automatically from the decaying step size—no multi-scale architecture or annealing schedule is required.

Table 2: **Ablation study.** MNIST ($\sigma=0.2$) PSNR (dB) at $K \in \{1, 5, 10, 20\}$. Condition A is the proposed KAN-EBM. Conditions B–D isolate architectural components. Conditions E–F vary filter-bank size.

	Condition	Params	$K=1$	$K=5$	$K=10$	$K=20$
A	KAN-EBM-full (ours)	31,782	22.92	23.91	25.12	27.21
B	MLP-EBM-filtered	70,801	22.66	22.66	22.65	22.65
C	KAN-EBM-raw (no filters)	2,336,000	23.70	27.12	28.47	27.21
D	FFN-DSM (strong baseline)	1,329,424	24.14	24.12	24.13	24.13
E	KAN-EBM-8filt	26,454	22.90	23.90	25.09	27.16
F	KAN-EBM-32filt	42,438	22.92	23.95	25.17	27.27

5.4 Ablation Study

The ablation study (Table 2) isolates four effects:

KAN vs. MLP (A vs. B). Replacing the KAN with an MLP energy (same filter bank, more parameters: 70K vs. 31K) eliminates test-time scaling entirely. MLP-EBM is flat at 22.65 ± 0.01 dB across all K . The *only* difference is the activation function: MLP uses fixed ReLU/GELU; KAN uses learnable B-splines. This confirms that learnable edge activations are the critical inductive bias.

Filter bank (A vs. C). Removing the filter bank (KAN-EBM-raw) increases parameters $74\times$ (31K $\rightarrow$ 2.3M). The raw KAN achieves a marginally better peak (28.5 dB at $K=10$) but *degrades* at $K=20$ (27.2 dB), suggesting instability at high K without the filter bank’s regularizing effect. The proposed KAN-EBM achieves comparable $K=20$ performance (27.2 dB) with $74\times$ fewer parameters.

Feedforward ceiling (A vs. D). FFN-DSM with 1.3M parameters achieves 24.1 dB in a single pass and cannot improve. KAN-EBM with 31K parameters surpasses this by 3.1 dB at $K=20$. The gain comes entirely from test-time compute, not additional capacity.

Filter count (A vs. E, F). 8, 16, and 32 filters yield 27.16, 27.21, and 27.27 dB at $K=20$ —a 0.11 dB range over a $5\times$ parameter variation. The filter bank is robust: 8 filters suffice.

5.5 Comparison to Diffusion and DEQ Baselines

Table 3: **K-scaling comparison at matched parameters (~ 31 – 50 K).** MNIST denoising ($\sigma=0.2$), PSNR (dB). KAN-EBM achieves 4.4 dB gain over $K=1$ – 20 ; DDPM-score gains 0.15 dB; DEQ gains 0.04 dB (converges in <5 steps by design).

Method	$K=1$	$K=2$	$K=5$	$K=10$	$K=20$
KAN-EBM (ours)	22.93	23.19	23.97	25.19	27.31
DDPM-score	22.66	22.67	22.70	22.74	22.81
DEQ	22.98	22.99	23.02	23.01	23.01

Table 3 compares K -step scaling across three methods at matched parameter counts (~ 31 – 50 K). KAN-EBM gains 4.4 dB from $K=1$ to $K=20$. DDPM-score, despite using the same Langevin-like update, gains only 0.15 dB because its explicit score network lacks a structured energy landscape to descend. DEQ gains only 0.04 dB: by design, fixed-point iteration converges in 1–2 steps, after which additional computation provides no benefit. The KAN’s energy landscape is uniquely suited for extended gradient descent.

5.6 Natural Image Denoising

Table 4 evaluates KAN-EBM on natural images. The model is trained on MNIST at $\sigma=0.2$ and *evaluated zero-shot* on CBSD68 and Set12—a significant domain shift. Despite this, it exhibits the same monotonic K -scaling on natural images. At $\sigma=15$, PSNR grows from 24.8 dB at $K=1$ to 28.1 dB at $K=20$ on CBSD68—a 3.3 dB gain, closing much of the gap to BM3D’s 33.5 dB (though our model was not trained on natural

Table 4: **Grayscale image denoising PSNR (dB)** on CBSD68 (Y-channel) and Set12. Classical and CNN baselines from published papers [Dabov et al., 2007, Zhang et al., 2017, 2018]. Our KAN-EBM shows K-scaling gain (numbers in parentheses show Δ from $K=1$). At $\sigma=15$, $K=20$ closes ≈ 3.3 dB of the gap to BM3D compared to $K=1$.

Dataset	Method	$\sigma = 15$			$\sigma = 25$			$\sigma = 50$		
		$K=1$	$K=10$	$K=20$	$K=1$	$K=10$	$K=20$	$K=1$	$K=10$	$K=20$
CBSD68	BM3D [Dabov et al., 2007]	33.52	—	—	30.71	—	—	27.38	—	—
	DnCNN [Zhang et al., 2017]	33.90	—	—	31.73	—	—	28.01	—	—
	FFDNet [Zhang et al., 2018]	33.87	—	—	31.63	—	—	28.05	—	—
	FFN	27.83	27.83	27.83	27.10	27.10	27.10	25.02	25.02	25.02
	MLP-EBM	25.49	25.69	25.66	20.90	21.30	21.16	15.08	15.73	14.78
	KAN-EBM (ours)	24.82	26.45 (+1.6)	28.14 (+3.3)	20.43	21.57 (+1.1)	22.82 (+2.4)	14.87	15.76 (+0.9)	16.74 (+1.9)
Set12	BM3D [Dabov et al., 2007]	33.06	—	—	31.42	—	—	28.83	—	—
	DnCNN [Zhang et al., 2017]	32.86	—	—	32.22	—	—	29.23	—	—
	FFN	27.58	27.58	27.58	26.69	26.69	26.69	24.20	24.20	24.20
	MLP-EBM	25.53	25.70	25.67	20.92	21.28	21.10	15.08	15.69	14.65
	KAN-EBM (ours)	24.84	26.52 (+1.7)	28.29 (+3.5)	20.45	21.59 (+1.1)	22.85 (+2.4)	14.87	15.75 (+0.9)	16.72 (+1.8)

images). The gain is consistent across all σ and both benchmarks, confirming that K-scaling generalizes beyond the training domain.

Note that our absolute numbers are below BM3D/DnCNN: these methods were designed and trained specifically for CBSD68, while KAN-EBM demonstrates a different capability (controllable scaling) rather than maximum absolute PSNR at $K=1$.

5.7 Interpretability Analysis

Spline activations. Figure 4 visualizes learned spline activations $\phi_{ij}(t)$ across input-to-hidden connections. The functions are diverse: some are monotonically increasing (edge detectors), others are peaked at interior points (band-pass responses), and a few are near-zero (unused connections). This diversity confirms that the KAN is exploiting its full representational capacity. Prior to adding the tanh squash and precision weighting, all activations collapsed near zero—a failure mode that motivated the architectural corrections described in section 3.1.

Energy landscape. Figure 5 shows slices of the learned energy landscape in a 2D subspace. The KAN carves deep, compact attractor basins around digit prototypes—exactly the geometry needed for reliable gradient descent. In contrast, MLP energy landscapes (not shown) are relatively flat with inconsistent gradient directions, explaining their failure to support iterative refinement.

Precision weights. Figure 6 shows the learned precision weights after training. The model assigns differentiated importance to filter channels: high-precision filters (corresponding to oriented edges and center-surround detectors, as seen in fig. 1) are amplified before the KAN; low-precision filters contribute minimally. This selective weighting is analogous to attention and is consistent with predictive coding’s precision-weighted error signals.

Energy convergence. Figure 7 confirms that the learned energy decreases monotonically during gradient descent for all test images. This is not guaranteed a priori: a poorly trained EBM can have energy landscapes where naive gradient descent oscillates or increases. The monotonic decrease validates that DSM training successfully shapes the gradient field into a reliable denoising direction.

Score alignment. Figure 8 measures how well the learned gradient field aligns with the ground-truth score direction. After 30 epochs of DSM training, the cosine similarity between $-\nabla E_\theta$ and the Tweedie-optimal direction exceeds 0.95, confirming that the KAN has learned a meaningful energy gradient, not an arbitrary one. This explains the PSNR improvements: each gradient step is a reliable stride toward the clean-image manifold.

5.8 Mathematical Validation

Table 5: **Mathematical validation.** Cosine similarities between gradient vectors of Allen-Cahn, score matching, and predictive coding. All three are numerically identical to floating-point precision. The symplectic vs. naïve integrator comparison demonstrates 7.3×10^8 better Hamiltonian conservation.

Experiment	Value	
$\cos(\text{Allen-Cahn, Score})$	1.00000000	✓
$\cos(\text{Allen-Cahn, PC-err})$	1.00000012	✓
$\cos(\text{Score, PC-err})$	1.00000012	✓
Symplectic max $ \Delta H $	1.37	
Naïve max $ \Delta H $	1.0×10^9	
Ratio (naïve / symplectic)	$7.3 \times 10^8 \times$	✓
PC free-energy decrease	100% monotone	✓

Table 5 reports the numerical validation of Section 4. The cosine similarities between the three inference algorithms are 1.0000 to floating-point precision, confirming exact analytical equivalence. Additionally, a symplectic integrator for the Hamiltonian dynamics conserves total energy to within 1.37 units across all steps, vs. 10^9 units for a naïve integrator—a $7.3 \times 10^8 \times$ improvement that validates the physical interpretation of the inference dynamics.

6 Related Work

Energy-based models. LeCun et al. [2006] introduced the EBM framework for perception and reasoning. Du and Mordatch [2019] showed that EBMs trained with MCMC can be used for image generation. Grathwohl et al. [2020] treated classifier logits as energies (JEM). Our work differs in using KAN parameterization and DSM training for *inference-time scaling*, not generation.

Score-based models. Song et al. [2021] unify score matching and diffusion through SDEs. Ho et al. [2020] demonstrate high-quality generation through DDPM chains. These models iterate at inference via a learned noise-conditioned score network; our KAN-EBM differs by using a single unconditional energy and gradient descent, producing a qualitatively different scaling behavior.

Kolmogorov-Arnold Networks. Liu et al. [2024] introduced KANs and demonstrated superior accuracy on scientific tasks. Follow-up work has applied KANs to physics simulation [Shukla et al., 2024] and time-series analysis. We are the first to use KANs as EBM energy parameterizations.

Deep Equilibrium Models. Bai et al. [2019] find implicit fixed points $z^* = f_\theta(z^*, \mathbf{u})$ via root-finding. Additional iterations do not improve quality once the fixed point is found (typically in 5–15 steps). This contrasts with KAN-EBM, where additional steps continue to improve by traversing the energy landscape.

Test-time compute scaling. Snell et al. [2024] show that allocating additional inference-time compute to language models via best-of-N sampling or tree search can match a $2 \times$ -larger model. Wu et al. [2024] extends this to diffusion models. KAN-EBM is a minimal, interpretable architecture that exhibits the same property for image reconstruction.

Predictive coding. Rao and Ballard [1999], Friston [2005], Bogacz [2017] formalize hierarchical perception as free-energy minimization. Our architecture inherits precision-weighted error minimization; we show numerically that our inference dynamics are identical to PC inference for Gaussian models.

7 Discussion

Why does the KAN enable scaling? We hypothesize that B-spline activations allow the KAN to represent energy landscapes with the following properties: (i) sharp, compact minima near data modes—so gradient descent converges quickly; (ii) smooth, informative gradients away from minima—so steps from noisy inputs always point toward the manifold. Fixed activations (ReLU, GELU) impose an algebraic structure on the

energy that prevents this: the piecewise-linear nature of ReLU creates flat regions with zero gradient where descent stalls.

The filter bank as inductive bias. The convolutional structure enforces translation equivariance and local feature extraction—priors well-matched to natural images. Without the filter bank (KAN-EBM-raw), the KAN must discover these spatial biases from data alone, requiring $74\times$ more parameters. The factored design is both efficient and interpretable.

Limitations. (i) Absolute PSNR at $K=1$ is below specialized denoisers (BM3D, DnCNN) trained on natural images. KAN-EBM’s advantage is controllable scaling with K , not static per-image quality. (ii) Training on MNIST and evaluating on CBSD68/Set12 involves domain shift; in-distribution training would improve absolute numbers. (iii) The model is currently limited to 28×28 patches; extension to high-resolution images requires hierarchical or patch-based processing.

Future directions. Natural extensions include: hierarchical KAN-EBMs with multiple spatial scales; learned step-size schedules (meta-learned η_t); and application to inverse problems beyond denoising (super-resolution, inpainting, compressed sensing). The free-energy connection to predictive coding suggests applying the framework to hierarchical latent variable models.

8 Conclusion

We introduced KAN-EBM, an architecture that pairs a convolutional filter bank with a Kolmogorov-Arnold Network to define a learnable, interpretable energy function. Trained with denoising score matching, KAN-EBM achieves monotonic PSNR improvement with inference steps: 4.3 dB gain from $K=1$ to $K=20$ on MNIST, 3.3 dB gain on natural images (CBSD68), and consistently flat or divergent baselines. Ablations confirm that the KAN’s learnable B-spline activations are the essential ingredient—replacing them with an MLP kills test-time scaling entirely. We further proved, and numerically validated, that our gradient-descent inference is mathematically identical to Allen-Cahn dynamics, score matching, and predictive-coding error minimization. These results position KAN-EBMs as a principled, efficient, and interpretable framework for controllable test-time compute in image reconstruction.

References

- S. Bai, J. Z. Kolter, and V. Koltun. Deep equilibrium models. In *Advances in Neural Information Processing Systems*, volume 32, 2019.
- R. Bogacz. A tutorial on the free-energy framework for modelling perception and learning. *Journal of Mathematical Psychology*, 76:198–211, 2017.
- B. Brown, J. Juravsky, R. Ehrlich, R. Clark, Q. V. Le, C. Ré, and A. Mirhoseini. Large language monkeys: Scaling inference compute with repeated sampling. *arXiv preprint arXiv:2407.21787*, 2024.
- K. Dabov, A. Foi, V. Katkovnik, and K. Egiazarian. BM3D image denoising with shape-adaptive principal component analysis. *Signal Processing with Adaptive Sparse Structured Representations (SPARS)*, 2007.
- Y. Du and I. Mordatch. Implicit generation and modeling with energy-based models. *Advances in Neural Information Processing Systems*, 32, 2019.
- K. Friston. A theory of cortical responses. *Philosophical Transactions of the Royal Society B: Biological Sciences*, 360(1456):815–836, 2005.
- W. Grathwohl, K.-C. Wang, J.-H. Jacobsen, D. Duvenaud, M. Norouzi, and K. Swersky. Your classifier is secretly an energy based model and you should treat it like one. In *International Conference on Learning Representations*, 2020.
- M. Gutmann and A. Hyvärinen. Noise-contrastive estimation: A new estimation principle for unnormalized statistical models. In *Proceedings of the Thirteenth International Conference on Artificial Intelligence and Statistics*, pages 297–304, 2010.

- G. E. Hinton. Training products of experts by minimizing contrastive divergence. volume 14, pages 1771–1800, 2002.
- J. Ho, A. Jain, and P. Abbeel. Denoising diffusion probabilistic models. In *Advances in Neural Information Processing Systems*, volume 33, pages 6840–6851, 2020.
- A. Hyvärinen. Estimation of non-normalized statistical models by score matching. *Journal of Machine Learning Research*, 6:695–709, 2005.
- Y. LeCun, S. Chopra, R. Hadsell, M. Ranzato, and F.-J. Huang. A tutorial on energy-based learning. In *Predicting Structured Outputs*. MIT Press, 2006.
- Z. Liu, Y. Wang, S. Vaidya, F. Ruehle, J. Halverson, M. Soljačić, T. Y. Hou, and M. Tegmark. KAN: Kolmogorov-Arnold networks. *arXiv preprint arXiv:2404.19756*, 2024.
- R. P. Rao and D. H. Ballard. Predictive coding in the visual cortex: a functional interpretation of some extra-classical receptive-field effects. *Nature Neuroscience*, 2(1):79–87, 1999.
- S. Roth and M. J. Black. Fields of experts. volume 82, pages 205–229, 2009.
- K. Shukla, J. D. Toscano, Z. Wang, Z. Zou, and G. E. Karniadakis. A comprehensive and FAIR comparison between MLP and KAN representations for differential equations and operator networks. *arXiv preprint arXiv:2406.02917*, 2024.
- C. Snell, J. Lee, K. Xu, and A. Kumar. Scaling LLM test-time compute optimally can be more effective than scaling model parameters. *arXiv preprint arXiv:2408.03314*, 2024.
- Y. Song, J. Sohl-Dickstein, D. P. Kingma, A. Kumar, S. Ermon, and B. Poole. Score-based generative modeling through stochastic differential equations. In *International Conference on Learning Representations*, 2021.
- P. Vincent. A connection between score matching and denoising autoencoders. *Neural Computation*, 23(7):1661–1674, 2011.
- N. Wu, Y. Hu, Y. J. Li, J. Sekhon, T. Jaakkola, H. He, and T. Salimans. Inference-time scaling for diffusion models beyond scaling denoising steps. *arXiv preprint arXiv:2501.09732*, 2024.
- K. Zhang, W. Zuo, Y. Chen, D. Meng, and L. Zhang. Beyond a Gaussian denoiser: Residual learning of deep CNN for image denoising. volume 26, pages 3142–3155, 2017.
- K. Zhang, W. Zuo, and L. Zhang. FFDNet: Toward a fast and flexible solution for CNN-based image denoising. volume 27, pages 4608–4622, 2018.




Figure 1: **Learned spatial filter bank.** Sixteen 5×5 convolutional filters learned from scratch via DSM training. Filters exhibit Gabor-like patterns, center-surround structures, and edge detectors at multiple orientations—reminiscent of V1 simple cells—demonstrating that the model discovers spatially interpretable features rather than arbitrary kernels. These features are passed through a tanh squash before entering the KAN.




Figure 2: **PSNR vs. inference steps K** on MNIST denoising ($\sigma=0.2$). KAN-EBM (blue) improves monotonically with K : +4.3 dB from $K=1$ to $K=20$. FFN-DSM (orange) is flat by construction—it has no iterative mechanism. MLP-EBM (red) diverges, indicating an ill-formed energy landscape. Shaded regions show ± 1 standard deviation across test images.




Figure 3: **Inference trajectory** for a single MNIST digit ($\sigma=0.2$). Left to right: noisy input, then KAN-EBM outputs at $K=1, 3, 5, 10, 20$, then ground truth. PSNR (dB) annotated above each frame. The denoised estimate progressively sharpens, recovering fine strokes by $K=10$ and achieving near-clean quality by $K=20$. Step size decays as $\eta_0 \cdot 0.97^t$, preventing oscillation at large K .




Figure 4: **Learned KAN spline activations** $\phi_{ij}(t)$ for a subset of input-to-hidden connections. Each panel shows one learned univariate function over the B-spline grid $[-1, 1]$. Functions exhibit diverse shapes: monotone, peaked, oscillatory, and near-zero. Non-trivial activations confirm that the spline basis is used across its full support, enabled by the tanh squash in section 3.1.




Figure 5: **Energy landscape slices.** The learned energy E_θ visualized in a 2D PCA-projected subspace of pixel space, centered on a test image. Clear attractor basins (blue minima) surround prototypical digit configurations. Noisy inputs (yellow points) sit outside these basins; gradient descent drives them inward (arrows), explaining why additional steps improve quality.




Figure 6: **Learned precision weights** $\Pi_j = \text{softplus}(\log \Pi_j)$ per filter channel after 30 epochs. Filters with high Π dominate the energy function; low- Π filters contribute weakly. The spread indicates the model has learned to prioritize certain spatial features (e.g., oriented edges) over others. Bars are sorted by Π value; filter index shown on x-axis.




Figure 7: **Energy** $E_\theta(u_t)$ **during inference** for 20 test images. All trajectories decrease monotonically (individual in gray; mean in blue), confirming that gradient descent reliably minimizes the learned energy. The convergence rate slows after $t \approx 10$ due to the decaying step size $\eta_t = 0.05 \cdot 0.97^t$, consistent with a well-formed energy basin.




Figure 8: **Score alignment.** Cosine similarity between $-\nabla_x E_\theta(\tilde{\mathbf{x}})$ (KAN-EBM score) and $(\mathbf{x} - \tilde{\mathbf{x}})/\sigma^2$ (ground-truth score direction) across 100 test images at $\sigma=0.2$. Mean cosine similarity > 0.95 at initialization, rising toward 1.0 after training. High alignment explains why gradient descent efficiently reduces denoising error.